

# Notes for Week 2 of CNN by Andrew. Ng @deeplearning.ai

---

Week 2 : Deep convolutional models: case studies

V1.0

17/02/2020

## Case Studies

---

### Why look at case studies?

- See someone others' codes can helps us understand better. We can perhaps use some ideas of others to build our own Neural Network.
- Some model of NN works well for certain project may work well for your projects, too.
- Help us to know how to read some **research papers**

### Classic Networks

#### LeNet-5

---

First introduced by [LeCun et al., 1998](#).

[LeCun et al., 1998. Gradient-based learning applied to document recognition]

#### Structure of Model

- input layer: digital picture of size  $(32 * 32 * 1)$
- 1st Convolutional layer: size  $\rightarrow (28 * 28 * 6)$ 
  - size =  $5 * 5$
  - stride = 1
  - num\_of\_channels = 6
- 1st AvePooling layer: size  $\rightarrow (14 * 14 * 6)$ 
  - $f = 2$
  - stride = 2
- 2nd Convolutional layer: size  $\rightarrow (10 * 10 * 16)$ 
  - size =  $5 * 5$
  - stride = 1
  - num\_of\_channels = 16
- 2nd AvePooling layer: size  $\rightarrow (5 * 5 * 16)$ 
  - $f = 2$
  - stride = 2
- 1st Full Connected layer: size =  $(5 * 5 * 16) = 400 \rightarrow 120$
- 2nd Full Connected layer: size =  $120 \rightarrow 84$
- output layer: output\_size = 1

## Pattern of Models

- As the network going deeper and deeper:
  - $n_H, n_W \downarrow$
  - $n_C \uparrow$
- Sequential:
  - Conv + Pool
  - Conv + Pool
  - Fc
  - Fc
  - ...
  - Output

## Guideline for reading the original paper

It's the most difficult paper introduced in this week, and it has used many techniques to reduce the computations. We will focus on the **part ii** and **part iii**.

## AlexNet

---

First introduced by [Alex Krizhevsky et al., 2012](#).

[Alex Krizhevsky et al., 2012. ImageNet classification with deep convolutional neural networks]

## Structure of Model

- input layer: digital picture of size  $(227 * 227 * 3)$
- 1st Convolutional layer: size  $\rightarrow (55 * 55 * 96)$ 
  - size =  $11 * 11$
  - stride = 4
  - num\_of\_channels = 96
- 1st MaxPooling layer: size  $\rightarrow (27 * 27 * 96)$ 
  - size =  $3 * 3$
  - stride = 2
- 1st **Same** Padding layer: size  $\rightarrow (27 * 27 * 256)$ 
  - size =  $5 * 5$
  - num\_of\_channels = 256
- 2nd MaxPooling layer: size  $\rightarrow (13 * 13 * 256)$ 
  - size =  $3 * 3$
  - stride = 2
- 2nd **Same** Padding layer: size  $\rightarrow (13 * 13 * 384)$ 
  - size =  $3 * 3$
  - num\_of\_channels = 384
- 3rd **Same** Padding layer: size  $\rightarrow (13 * 13 * 384)$ 
  - size =  $3 * 3$
  - num\_of\_channels = 384
- 4th **Same** Padding layer: size  $\rightarrow (13 * 13 * 256)$

- size =  $3 * 3$
- num\_of\_channels = 256
- 3rd MaxPooling layer: size  $\rightarrow (6 * 6 * 256) = 9216$ 
  - size =  $3 * 3$
  - stride = 2
- 1st Full Connected layer: size =  $9216 \rightarrow 4096$
- 2nd Full Connected layer: size =  $4096 \rightarrow 4096$
- output layer: size\_of\_Softmax = 1000

### Patterns of Model & Paper

- Similarity to [LeNet](#), but much more bigger
- Use Relu as an activation function
- Mentioned for the [paper](#):
  - Multiple GPUs for Conv layers to gain a higher speed
  - Local Response Normalization(LRN)
    - Normalize at each position of the layer with depth of all the chanel.
  - *Easiest paper to read introduced this week*

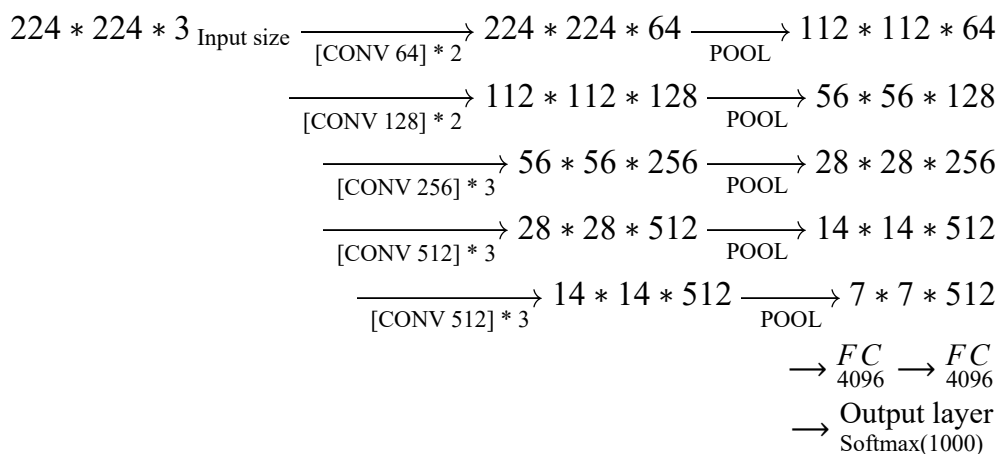
### VGG-16

Introduced by [Simonyan&Zisserman 2015](#). Using much more simple layers

- CONV :  $3 * 3$  filter,  $s = 1$ , padding = same
- MAX-POOL :  $2 * 2$   $s = 2$

### Structure of Model

With [CONV] and [POOL] layers defined above:



### Patterns of Model

- It's really a huge Neural Network
  - ~ 138M variables
- architecture uniform
- roughly double num of channels and divide with pooling
  - $n_H, n_W \downarrow \quad * = \frac{1}{2}$

- $n_C \uparrow$   $\quad \quad \quad * = 2$
- which is very interesting to reserve this relationship

## Residual Network

Introduced by [He et al., 2015](#)

He et al., 2015. Deep residual networks for image recognition

- Very, very deep neural networks are difficult to train because of vanishing and exploding gradient types of problems.
- Residual block
  - skip connections which allows you to take the **activation** from one layer and **suddenly** feed it to another layer even much deeper in the neural network
  - after the **linear** part but before the **activation** part
  - use a *skip cut* or "skip connection"
  - $a^{[l+2]} = g(z^{[l+2]} + a^{[l]})$
- in theory deep nn gives better result but in reality no
- resNet can give better result with deep NN

## Why ResNets Works?

- classically, when neural network goes deeper and deeper its ability will go down in contrast to our wishes.
- after big nn, we use residual block, which:
  - at least do no harm to the forward deep NN:  
We use **ReLU** here as activation function, so there will be no-negative activation.  
As we know that the residual block works as

$$\begin{aligned} a^{[l+2]} &= g(z^{[l+2]} + a^{[l]}) \\ &= g(w^{[l+2]}a^{[l+1]} + b^{[l+1]} + a^{[l]}) \end{aligned}$$

**IF** there is nothing to learn, then suppose that

$$w^{[l+2]} = b^{[l+1]} = 0$$

we'll still get

$$\begin{aligned} a^{[l+2]} &= g(z^{[l+2]} + a^{[l]}) \\ &= g(w^{[l+2]}a^{[l+1]} + b^{[l+1]} + a^{[l]}) \\ &\xrightarrow{w^{[l+2]}=b^{[l+1]}=0} g(a^{[l]}) \\ &= a^{[l]} \end{aligned}$$

**ELSE** we can perhaps learn something in these residual layers

- In fact, in traditional deep network, those deep layers have huge difficulty to choose parameters in order to get better performance

## Networks in Networks and 1\*1 Convolutions

2020-02-17

First introduced by [Lin et al., 2013. Network in Network](#) what does a  $1 \times 1$  convolution do?

- multiply?
- In the case of multiple channels,
  - in reality, these namely  $1 \times 1$  Convolution layers have the dimension of  $(1 \times 1 \times n_C^{[l-1]})$ .
  - Thus, if we have the number of  $1 \times 1$  filters is  $n_C^{[l]}$ , we will get the layer  $[l]$  of original height & weight dimensions and new channel dimension.
- depend on channels + relu
- change num of channels
  - shrink
  - keep
  - or even increase
  - *by modifying num of filters of these  $1 \times 1$  Convolutions*

## Inception Network

- why should i chose all those size of conv pool...
- First introduced by [Szegedy et al., 2014. Going deeper with convolutions](#)

### Motivation for inception network

- stack the results of different conv layers (Padding = Same)
    - $1 \times 1$  convolution
    - $3 \times 3$  same convolution
    - $5 \times 5$  same convolution
    - same Max-Pooling
    - ...
- These will result in different kind of channels with the **same size**

So that we can easily **stack them up** to form a new layer of multiple channels come from different **convolutions** and **poolings**

- in this way we do not need to choose size of convolutions
- Problem of computational cost
  - ex.  $5 \times 5$  filters

$$28 \times 28 \times 192 \xrightarrow[\text{same, 32}]{\text{CONV } 5 \times 5} 28 \times 28 \times 32$$

Will takes

$$28 * 28 * 32 * 5 * 5 * 192 \simeq 120M \text{ computations}$$

It's very expensive in fact

- in revanch, if we use a  $1 \times 1$  to first shrink the size to

$$28 \times 28 \times 16$$

can definitely reduce this cost:

$$28 \times 28 \times 192 \xrightarrow[16]{\text{CONV } 1 \times 1} 28 \times 28 \times 16 \xrightarrow[\text{same, 32}]{\text{CONV } 5 \times 5} 28 \times 28 \times 32$$

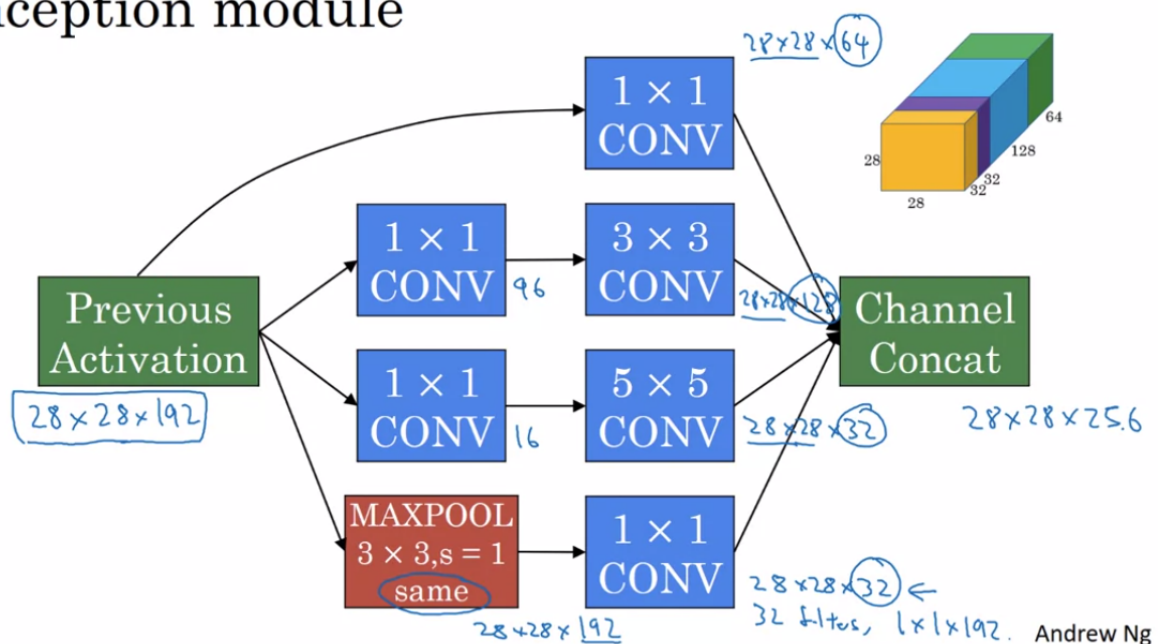
Will takes

$$28 * 28 * 16 \times 1 * 1 * 192 + 28 * 28 * 32 \times 5 * 5 * 16 \\ \simeq 2.4M + 10.0M = 12.4M$$

computations

Build an Inception Module

## Inception module



- In the paper:
  - what do side branches do?  
Side branches just try to predict the result as the output layer
  - first developed by google  
so we call it **goLeNet**
  - Where comes from the name Inception:  
**we need to go deeper**  
from film **Inception**

## Practical advices for using ConvNets

### Using Open-Source Implementation

- Hard to replicate the result of jobs of papers
- Online open source!
- Find them in Github
- Advantages:
  - Normally good neural networks take long time and many GPUs to train. So we can take advantage of results of those open source to do **transfer learning**

- We can also **Contribute** to the world
- Much more easy to **Get start!**

## Transfer Learning

- Someone else has already trained very deep neural networks which may take some weeks and also many GPUs computation
- Transfer knowledge is very useful here!

### How to do it?

- download some network implementation

not only the **code**

but also the **parameters/weights** that have been beautifully trained

- get rid of some final output decision layers and use our own layers to give an output which corresponds to our own problems
- Freeze the forward layers by setting:
  - trainable parameter = 0
  - freeze = 1that is to say, these forward layers will be **frozen** so that we **don't** train their parameters but just use the result of ancient training.

### We can also:

- **map** the input to the **last** frozen layer  
because these layers are already trained, we just do some **pre-compute** so that we can save the result to disk for later training.
- freeze fewer layers when we have a big dataset to learn from
  - we can just add more **our own hidden layers**
  - General pattern:  
more data --> fewer layers frozen

## Data Augmentation

- Computer vision is very complex task so we need as much as possible data
- Often data augmentation will help

### Common augmentation method

1. Mirroring
2. Random Cropping
3. Rotation
4. Shearing
5. Local warping
6. ...

### Color shifting

- $R += 20$
- $G -= 20$
- $B += 20$   
More robust to the color changes
- PCA(Principle Component Analysis)
- [ml-class.org](http://ml-class.org)
- AlexNet paper PCA color augmentation

## Implementing distortions during training

- one thread and other thread
- CPU thread works to form a mini-batch of data
- While GPU or other CPU will work for other training procedure
- These can be done in parallel

## State of Computer Vision

### Data vs. hand-engineering

larger and larger dataset

- object detection
- Image recognition
- Speech recognition

In history, when we do not have so much data, people tends to use a lot of **Hacks** (more **hand-engineering**). But now since there's larger and larger dataset, people tends to use some very simple architecture of algorithm.

Two sources of knowledge

- Labeled data
- hand engineered features/network architecture/other components

Often we donnot have as much data as we want

- when we have little data, there will be more techniques

### Tips fpr doing well on benchmarks/winning comprtitions

Ensembling

- Train several networks independently and average their outputs  
outputs but not the weights

Multi-crop at test time

- Run classifier on multiple versions of test images and average results



## Use open source code

- Use architecture of networks published in the literature
- Use open source implementations if possible
- Use pretrained models and fine-tune on your dataset